Uppsala University

Bayesian Statistics and Data Analysis

Assignment 7

General information

- The recommended tool in this course is R (with the IDE R-Studio). You can download R **here** and R-Studio **here**. There are many tutorials, videos and introductions to R and R-Studio online. You can find some initial hints from **RStudio Education pages**.

- When working with R, we recommend writing the report using R markdown and the provided **R markdown template**. The template includes the formatting instructions and how to include code and figures.

- Instead of R markdown, you can use other software to make the PDF report, but you should use the same instructions for formatting. These instructions are also available in **the PDF produced from the R markdown template**.

- We supply a Google Colab notebook that you can also use for the assignments. We have included the installation of all necessary R packages; hence, this can be an alternative to using your own local computer. You can find the notebook **here**. You can also open the notebook in Colab **here**.

- Report all results in a single and *anonymous* pdf. Note that no other formats are allowed.

- The course has its own R package `bsda` with data and functionality to simplify coding. To install the package, just run the following (upgrade="never" skips question about updating other packages):

  1. `install.packages("remotes")`
  2. `remotes::install_github("MansMeg/BSDA",`
     `subdir = "rpackage", upgrade="never")`

- Many of the exercises can be checked automatically using the R package `markmyassignment`. you can find information on how to install and use the package **here**. There is no need to include `markmyassignment` results in the report.

- You can find common questions and answers regarding the installation and technical problems in Frequently Asked Questions (FAQ).

- You can find deadlines and information on how to turn in the assignments in Studium.

- You are allowed to discuss assignments with your friends, but it is not permitted to copy solutions directly from other students or the internet. Try to solve the actual assignment problems with your code and explanations. Do not share your answers publicly. We compare the answers with the "urkund" system. We will report all suspected plagiarism.

- If you have any suggestions or improvements to the course material, please post in the course chat feedback channel, create an issue, or submit a pull request to the public repository **here**.

- It is *mandatory* to include the following parts in all assignments (these are included already in the template):

  1. Time used for reading: How long time took the reading assignment (in hours)
  2. Time used for the assignment: How long time took the basic assignment (in hours)
  3. Good with assignment: Write one-two sentences of what you liked with the assignment/what we should keep for next year.
  4. Things to improve in the assignment: Write one-two sentences of what you think can be improved in the assignment. Can something be clarified further? Did you get stuck on stuff unrelated to the content of the assignment etc.

- You can find information on how each assignment will be graded and how points are assigned **here**. **Note!** This grading information can change during the course, for example, if we find errors or inconsistencies. Please feel free to comment on these grading instructions, ideally before turning in your assignment, if you think something is missing or is incorrect.

- To pass (G) the assignment, you need 70% of the total points. To pass with distinction (VG), you need 90% of the total points. See the grading information on the point allocations for each assignment.

Information on this assignment

This assignment is related to Chapter 5.

Reading instructions: Chapter 5 in BDA3, see reading instructions.

Reporting accuracy: For posterior statistics of interest, only report digits for which the Monte Carlo standard error (MCSE) is zero. *Example:* If you estimate $E(\mu) = 1.234$ with $\text{MCSE}(E(\mu)) = 0.01$, you should report $E(\mu) = 1.2$.

Installing and using stan: To install Stan on your laptop, https://github.com/stan-dev/rstan/wiki/RStan-Getting-Started. If you encounter problems, see additional answers in the **FAQ**. Recently there have been reports of installation problems with Windows and R 4.0 (see Stan discourse for more).

Installing and using CmdStanR: If you want to use Stan in R on local computer, it can be easier to install CmdStanR interface mc-stan.org/cmdstanr/.

Additional useful packages are loo, bayesplot and shinystan (but you don't need these in this assignment). For Python users, PyStan, CmdStanPy, and Arviz packages are useful.

Stan manual can be found at https://mc-stan.org/users/documentation/. From this website, you can also find a lot of other useful material about Stan.

1. Linear model: drowning data with Stan

The provided data `drowning` in the `bsda` package contains the number of people who died from drowning each year in Finland 1980–2019. A statistician is going to fit a linear model with Gaussian residual model to these data using time as the predictor and number of drownings as the target variable (see the related linear model example for the Kilpisjärvi-temperature data in the example Stan codes). She has two objective questions:

i) What is the trend of the number of people drowning per year? (We would plot the histogram of the slope of the linear model.)

ii) What is the prediction for the year 2020? (We would plot the histogram of the posterior predictive distribution for the number of people drowning at $\tilde{x} = 2020$.)

To access the data, use:

```
library(bsda)
data("drowning")
```

Corresponding Stan code is provided in Listing 1. However, it is not entirely correct for the problem. First, there are *three mistakes*. Second, there are no priors defined for the parameters. In Stan, this corresponds to using uniform priors.

Your tasks are the following:

a) Find the three mistakes in the code and fix them. Report the original mistakes and your fixes clearly in your report. Include the *full* corrected Stan code in your report.

   Hint: You may find some of the mistakes in the code using Stan syntax checker. If you copy the Stan code to a file ending .stan and open it in RStudio (you can also choose from RStudio menu File→New File→Stan file to create a new Stan file), the editor will show you some syntax errors. More syntax errors might be detected by clicking 'Check' in the bar just above the Stan file in the RStudio editor. Note that some of the errors in the presented Stan code may not be syntax errors.

b) Determine a suitable weakly-informative prior normal$(0, \sigma_\beta)$ for the slope `beta`. It is very unlikely that the mean number of drownings changes more than 50 % in one year. The approximate historical mean yearly number of drownings is 138. Hence, set σ_β so that the following holds for the prior probability for `beta`: $\Pr(-69 < \text{beta} < 69) = 0.99$. Determine suitable value for σ_β and report the approximate numerical value for it.

c) Using the obtained σ_β, add the desired prior in the Stan code. In the report, in a separate section, indicate clearly how you carried out your prior implementation, e.g. "Added line ... in block ...".

d) In a similar way, add a weakly informative prior for the intercept `alpha` and explain how you chose the prior.

Hint! Example resulting plots for the problem, with the fixes and the desired prior applied, are shown in Figure 1. If you want, you can use these plots as a reference for testing

4

if your modified Stan code produces similar results. However, running the inference and comparing the plots is not required.

Note! The example/test plots and results are based on data up to 2016. You should report your result for the whole period 2019.
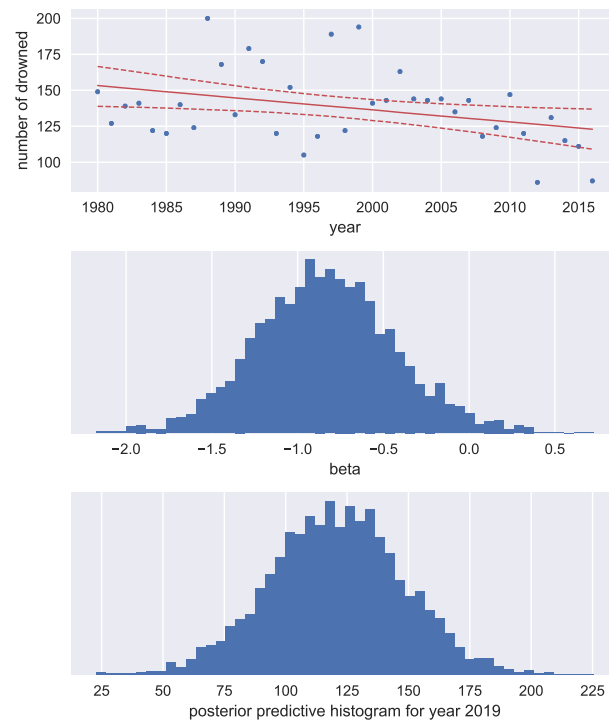

Figure 1: Example plots for the results obtained for the problem in the Question 1 with data until 2016. In the first subplot, the red lines indicate the resulting 5 %, 50 %, and 95 % posterior quantiles for the transformed parameter `mu` at each year.

Listing 1: Broken Stan code for question 1

```
 1  data {
 2      int<lower=0> N;   // number of data points
 3      vector[N] x;      // observation year
 4      vector[N] y;      // observation number of drowned
 5      real xpred;       // prediction year
 6  }
 7  parameters {
 8      real alpha;
 9      real beta;
10      real<upper=0> sigma;
11  }
12  transformed parameters {
13      vector[N] mu = alpha + beta*x;
14  }
15  model {
16      y ~ normal(mu, sigma)
17  }
18  generated quantities {
19      real ypred = normal_rng(mu, sigma);
20  }
```

2. Hierarchical model: factory data with Stan

Note! Assignment 8 build upon this part of the assignment, so it is important to get this assignment correct before you start with Assignment 8.

The `factory` data in the `bsda` package contains quality control measurements from 6 machines in a factory. In the data file, each column contains the measurements for a single machine. Quality control measurements are expensive and time-consuming, so only 5 measurements were done for each machine. In addition to the existing machines, we are interested in the quality of another machine (the seventh machine). To read in the data, just use:

```
library(bsda)
data("factory")
```

For this problem, you'll use the following Gaussian models:

- a separate model, in which each machine has its own model/parameters

- a pooled model, in which all measurements are combined and there is no distinction between machines

- a hierarchical model, which has a hierarchical structure as described in BDA3 Section 11.6.

As in the model described in the book, use the same measurement standard deviation σ for all the groups in the hierarchical model. In the separate model, however, use separate measurement standard deviation σ_j for each group j. You should use weakly informative priors for all your models.

The provided Stan code in Listing 2 given on the next page is an example of the separate model (but with very strange results, why?). This separate model can be summarized mathematically as:

$$y_{ij} \sim N(\mu_j, \sigma_j)$$
$$\mu_j \sim N(0, 1)$$
$$\sigma_j \sim \text{Inv-}\chi^2(10)$$

To run Stan for that model, simply use:

```
data("factory")

sm <- rstan::stan_model(file = "[path to stan model code]")

stan_data <- list(
  y = factory,
  N = nrow(factory),
  J = ncol(factory)
)
model <- rstan::sampling(sm, data = stan_data)
model
```

```
## Inference for Stan model: 5cbfa723dd8fb382e0b647b3943db079.
## 4 chains, each with iter=2000; warmup=1000; thin=1;
## post-warmup draws per chain=1000, total post-warmup draws=4000.
##              mean se_mean   sd     2.5%     25%      50%
## mu[1]        0.11    0.01  0.98    -1.81   -0.56    0.12    0.77
## mu[2]        0.10    0.01  1.00    -1.86   -0.56    0.10    0.79
## ...
```

Note! These are *not* the results you would expect to turn in your report. You will need to change the code for the separate model as well.

For *each of the three models* (separate, pooled, hierarchical), your tasks are the following:

a) Describe the model with mathematical notation (as is done for the separate model above). Also describe in words the difference between the three models.

b) Implement the model in Stan and include the code in the report. Use weakly informative priors for all your models.

c) Using the model (with weakly informative priors) report, comment on and, if applicable, plot histograms for the following distributions:

   i) the posterior distribution of the mean of the quality measurements of the sixth machine.

   ii) the predictive distribution for another quality measurement of the sixth machine.

   iii) the posterior distribution of the mean of the quality measurements of a seventh machine (not in the data).

d) Report the posterior expectation for μ_1 with a 90% credible interval but using a normal$(0, 10)$ prior for the μ parameter(s) and a Gamma$(1, 1)$ prior for the σ parameter(s). For the hierarchical model, use the normal$(0, 10)$ and Gamma$(1, 1)$ as hyper-priors.

8

Listing 2: Stan code for a bad separate model

```
1  data {
2    int<lower=0> N;
3    int<lower=0> J;
4    vector[J] y[N];
5  }
6
7  parameters {
8    vector[J] mu;
9    vector<lower=0>[J] sigma;
10 }
11
12 model {
13   // priors
14   for (j in 1:J){
15     mu[j] ~ normal(0, 1);
16     sigma[j] ~ inv_chi_square(10);
17   }
18
19   // likelihood
20   for (j in 1:J)
21     y[,j] ~ normal(mu[j], sigma[j]);
22 }
23
24 generated quantities {
25   real ypred;
26   // Compute predictive distribution
27   // for the first machine
28   ypred = normal_rng(mu[1], sigma[1]);
29 }
```